

COL362/632 Introduction to Database Management Systems

Recovery

Kaustubh Beedkar


Department of Computer Science and Engineering
Indian Institute of Technology Delhi



Admistrivia


- ▶ MyAwesomeDB portal will re-open tomorrow for the leaderboard task
 - ▶ Same DB schema but with higher scale factor (min SF 1)
 - ▶ Initial leaderboard will be based on your last submission (before 15 EOD)
 - ▶ Can (re)submit until 27th EOD
-
- ▶ Quiz-06 on Monday
 - ▶ Quiz-07 on Monday, depending on today's class

Introduction



Recall Buffer Pool Manager


- ▶ We use Buffer Pool Manager to read and write to disk
- ▶ Actual writing to disk is dictated by the buffer replacement policy



Failure Scenario I

- ▶ Consider the following schedule

T_1
—
read(A)
A:=A-50
write(A)
...
commit




- ▶ Transaction committed before page was flushed to disk—(Durability challenges!)

Failure Scenario II

- Consider the following schedule

T_1	T_2
read(A)	
write(A)	
	read(B)
	write(B)
	commit
...	
abort	



- Changes of uncommitted Tx flushed to disk—(Atomicity challenges!)

Why do Transactions Fail?

1. Transaction Failures

- Logical Errors (e.g., some integrity constraint violation)
- Internal State Errors (e.g., deadlock)

2. System Failure

- Software failures – OS or DBMS implementation (e.g., uncaught divide by zero)
- Hardware failures (e.g., power goes off)
- **Fail-stop assumption** – non-volatile storage contents are assumed to not be corrupted by system crash

3. Storage Media Failures

- Disk failure (e.g., a disk head crash)
 - destruction can be detected (use checksums to detect failures)
- Note: the database cannot recover from this! Restore from an archived version

DBMS must ensure

ACID: Changes from any Tx are durable once it has been committed

ACID: No partial changes are durable if the Tx are aborted

Logging & recovery

- ▶ Actions taken during normal execution to ensure that DBMS can recover from failure
- ▶ Actions taken after a failure to recover the database to a state that ensures consistency, atomicity, and durability

Redo & Undo

The two issues

Changes from any Tx are durable once it has been committed

- ▶ **Redo** certain actions that were committed but not written to disk

No partial changes are durable if the Tx are aborted

- ▶ **Undo** certain actions that did not commit

Storage Types

► Volatile Storage

- Does not survive system crash
- e.g., DRAM, cache memory


► Non-volatile Storage

- Survives system crash
- e.g., HDD, SSD

► Stable Storage

- A mythical form of storage that survives all possible failures
- Approximated by replication

Storage hierarchy



Some Terminology

- ▶ $\text{read}(A)$
- ▶ $\text{write}(A)$
- ▶ $\text{input}(B_A)$
 - Transfer the physical block B_A where the data items A resides to buffer pool
- ▶ $\text{output}(B_A)$
 - Transfer the buffer pool page B_A that had the data item A to disk (replace the old block B_A on disk)
 - Note: $\text{output}(B_A)$ may not happen immediately following $\text{write}(A)$

In the context of transactions

- ▶ A Tx must perform $\text{read}(A)$ before accessing A for the first time
- ▶ $\text{write}(A)$ can be executed anytime before the transaction commits

Log-based Recovery

Key Idea

Output first the information describing the modifications to stable storage without modifying the database

Log File

- ▶ A log file contains the changes
- ▶ When a Tx T_i starts, it registers itself by writing a $\langle T_i, start \rangle$ log record
- ▶ Before T_i executes a write(A), a log record $\langle T_i, A, V_{old}, V_{new} \rangle$ is written
- ▶ When T_i finishes its last statement, the log record $\langle T_i, commit \rangle$ is written

Logging allows to perform both undo and redo operations

Immediate-modifications

- ▶ Perform updates to buffer pool/disk before the Tx commits

Deferred-modifications

- ▶ Perform updates to buffer pool/disk only at the time of Tx commit

Note: All modification must be preceded by creation of a log record

Concurrency & Recovery

- ▶ Consider this schedule

T_1	T_2
...	
write(A)	
	...
	write(A)
abort	
	abort

- ▶ If T_1 aborts, we undo(T_1), and also have to undo(T_2)
- ▶ Require that if data item has been modified by a transaction, no other transaction can modify the data item until the first transaction commits or aborts
 - automatically happens in 2PL
 - time-stamp ordered protocol (recall dealing with cascaded aborts)
 - validation-based also supports this

When does a Tx go in committed state?

- ▶ when the commit log record $\langle T_i, commit \rangle$ (last record of a Tx) has been written to stable storage
- ▶ we don't care if the database has been modified or not

Also, recall partially committed state

Undo and Redo Operation

redo(T_i)

- ▶ Sets the value of all data items updated by T_i to the new values, going from first log record for T_i
- ▶ No logging is done in this case

undo(T_i)

- ▶ Restore the value of all data items updated by T_i to their old value, going backwards from the last record of T_i
 - Also write a **redo-only** record $\langle T_i, A, V_{old} \rangle$
 - When an undo(T_i) finishes, a $\langle T_i, abort \rangle$ log record is written out

When recovering from failure

- ▶ T_i needs to be undone if the log
 - contains the record $\langle T_i, start \rangle$
 - but does not contain either the record $\langle T_i, commit \rangle$ or $\langle T_i, abort \rangle$
- ▶ T_i needs to be redone if the log
 - contains the records $\langle T_i, start \rangle$
 - and either the record $\langle T_i, commit \rangle$ or $\langle T_i, abort \rangle$

Example

Consider the schedule

T_1	T_2
begin	
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
write(B)	
	begin
	read(C)
	$C := C - 100$
	write(C)

Scenario-1: System crash after write(B)

Log records

$\langle T_1 \text{start} \rangle$

$\langle T_1, A, 100, 50 \rangle$

$\langle T_1, B, 100, 150 \rangle$

Recovery action:

- ▶ $\text{undo}(T_1)$: B is restored to 100; A to 100
- ▶ and write log records
 $\langle T_1, B, 100 \rangle, \langle T_1, A, 100 \rangle, \langle T_1, \text{abort} \rangle$

Example

Consider the schedule

T_1	T_2
begin	
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
write(B)	
commit	
	begin
	read(C)
	$C := C - 100$
	write(C)

Scenario-2: System crash after write(C)

Log records

$\langle T_1, start \rangle$
 $\langle T_1, A, 100, 50 \rangle$
 $\langle T_1, B, 100, 150 \rangle$
 $\langle T_1, commit \rangle$
 $\langle T_2, start \rangle$
 $\langle T_2, C, 150, 50 \rangle$

Recovery action:

- ▶ redo(T_1) and undo(T_2): A is set to 50; B to 150, and C is restored to 150
- ▶ and write log records
 $\langle T_2, C, 150 \rangle, \langle T_2, abort \rangle$

Example

Consider the schedule

T_1	T_2
begin	
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
write(B)	
commit	
	begin
	read(C)
	$C := C - 100$
	write(C)
	commit

Scenario-3: System crash after T_2 commits

Log records

$\langle T_1, start \rangle$
 $\langle T_1, A, 100, 50 \rangle$
 $\langle T_1, B, 100, 150 \rangle$
 $\langle T_1, commit \rangle$
 $\langle T_2, start \rangle$
 $\langle T_2, C, 150, 50 \rangle$
 $\langle T_2, commit \rangle$

Recovery action:

- ▶ redo(T_1) and redo(T_2): A is set to 50; B to 150. Then C is set to 50

Checkpoints

- ▶ Redo and Undo operations can have significant overhead
 - Processing entire log can be time consuming
 - Redoing Tx that have already written to disk is wasteful
 - ▶ **Checkpoints**
 - Special marker: recovery only needs to look at parts of the log after checkpoint and just before the checkpoint
 - ▶ Checkpointing is done periodically
 - output all log records from buffer pool to stable storage
 - output all modified buffer blocks to disk
 - Write a log record $\langle \textit{checkpoint}, L \rangle$, where L is a list of all transactions active at the time of checkpoint
- No Tx allowed to update while checkpointing is in progress

Write-Ahead Logging

Key Ideas

- ▶ Tx T_i enters commit state only after $\langle T_i, commit \rangle$ log record has been output to disk
- ▶ Before the $\langle T_i, commit \rangle$ log record can be output to stable storage, all log records must have been output to stable storage
- ▶ Write to disk the log file records corresponding to database modifications **before** the buffer pool manager flushes pages to disk
- ▶ Buffer Pool Policy: Steal + No-force
 - Steal policy: Pages of uncommitted Tx may be written to disk
 - No-force policy: Allow Tx to commit even if it has modified pages that have not been written to disk

Write-Ahead Logging

Consider the schedule

T_1

begin

read(A)

write(A)


read(B)

write(B)

...

commit


- ▶ Assume that the modified page (with $A = 50$) had to be replaced
- ▶ First write the log page to disk, then flush the page




Write-Ahead Logging (scenario 2)

Consider the schedule

T_1
begin
read(A)
write(A)
read(B)
write(B)
...
commit



- ▶ Assume that T_1 is ready to commit
- ▶ First write the log page to disk, then commit the Tx



Optimization

- ▶ Group commits: batch multiple commits

Recovery Algorithm (sketch)

Write ahead logging (during normal execution)

- ▶ $\langle T_i \text{ start} \rangle$ when the Tx begins
- ▶ $\langle T_i, A, V_{old}, V_{new} \rangle$ for each write
- ▶ $\langle T_i \text{ commit} \rangle$ when Tx ends

Recovery Algorithm (sketch)

In case of transaction failure

assuming T_i needs to be rolled back

- ▶ Scan log backwards
- ▶ For each log record of the form $\langle T_i, X, V_{old}, V_{new} \rangle$
 - perform undo by writing V_{old} to X
 - write a **compensation log record** $\langle T_i, X, V_{old} \rangle$
- ▶ After $\langle T_i \text{ start} \rangle$ is found, stop scanning and write log record $\langle T_i \text{ abort} \rangle$

Recovery Algorithm (sketch)

In case of system failure

1. Redo phase: replay updates of **all** transactions

- ▶ Find last checkpoint $\langle \text{checkpoint } L \rangle$ record and set undo-list to L
- ▶ Scan log forward from $\langle \text{checkpoint } L \rangle$ record
 - whenever a record $\langle T_i, X, V_{old}, V_{new} \rangle$ is found, redo by writing V_{new} to X
 - whenever a record $\langle T_i \text{ start} \rangle$ is found, add T_i to undo-list
 - whenever a record $\langle T_i \text{ commit} \rangle$ or $\langle T_i \text{ abort} \rangle$ is found, remove T_i from undo-list

Recovery Algorithm (sketch)

In case of system failure

2. Undo phase:

- ▶ Scan log backwards from end
 - whenever a record $\langle T_i, X, V_{old}, V_{new} \rangle$ is found, where T_i is in undo list
 - perform undo by writing V_{old} to X
 - write a **compensation log record** $\langle T_i, X, V_{old} \rangle$
 - whenever a record $\langle T_i \text{ start} \rangle$ is found where T_i is in undo list
 - write log record $\langle T_i \text{ abort} \rangle$
 - remove T_i from undo-list
 - stop when undo-list is empty

Recovery Example

- ▶ Consider the following log

$\langle T_1 \text{ start} \rangle$
 $\langle T_1, B, 100, 150 \rangle$
 $\langle T_2 \text{ start} \rangle$
 $\langle \text{checkpoint } \{ T_1, T_2 \} \rangle$
 $\langle T_2, C, 100, 200 \rangle$
 $\langle T_2 \text{ commit} \rangle$
 $\langle T_3 \text{ start} \rangle$
 $\langle T_3, A, 100, 50 \rangle$
 $\langle T_1, B, 100 \rangle$
 $\langle T_1 \text{ abort} \rangle$
 $\langle T_3, A, 100 \rangle$
 $\langle T_3 \text{ abort} \rangle$

Scan forward from checkpoint (redo phase)

- ▶ Undo-list: T_1 ~~T_1~~ , T_2 ~~T_2~~ , T_3
- ▶ After $\langle T_2 \text{ commit} \rangle$ Remove T_2 from undo-list
- ▶ After $\langle T_3 \text{ start} \rangle$ Add T_3 to undo-list
- ▶ After $\langle T_1 \text{ abort} \rangle$ Remove T_1 from undo-list

Scan backward (undo phase)

- ▶ After writing 100 to A, Write compensation log record $\langle T_3, A, 100 \rangle$
- ▶ After $\langle T_3 \text{ start} \rangle$, write $\langle T_3 \text{ abort} \rangle$